

---

curvefs 2.4.0版本总体测试

- 一、遗留问题列表
- 二、测试内容和结论概述
- 三、测试要点
- 四、测试结论
- 五、详细测试数据及监控数据
  - 5.1 常规测试
    - 5.1.1 文件系统POSIX 接口
      - 5.1.1.1 pjdtest
      - 5.1.1.2 ltp-fsstress
        - 测试结果:
    - 5.1.2 元数据项& 数据属性
      - 5.1.2.1 dbench
      - 5.1.2.2 iozone
      - 5.1.2.4 rename 测试用例集
      - 5.1.2.5 xfstest
    - 5.1.3 数据一致性测试
      - 5.1.3.1 编译项目或者内核
      - 5.1.3.2 vdbench读写一致性测试
  - 5.2 异常测试
  - 5.3 新增功能测试
    - 5.3.1 warmup测试
      - 5.3.1.1 cto open
        - 5.3.1.1.1 静态warmup
        - 5.3.1.1.2 同时有读写时warmup
          - 5.3.1.1.2.1 缓存盘容量不足时
            - 5.3.1.1.2.1.1 大文件（根据缓存盘容量）并发操作
            - 5.3.1.1.2.1.2 大规模目录（1000w+）
          - 5.3.1.1.2.2 缓存盘容量足够时
            - 5.3.1.1.2.2.1 大文件（根据缓存盘容量）并发操作
            - 5.3.1.1.2.2.2 大规模目录（1000w+）
  - 5.4 回归测试
  - 5.5 2.3升级2.4 测试

## 一、遗留问题列表

问题列表

	风险项	ISSUE. No	负责人	严重级别	是否解决	是否需要回归	回归人	是否回归通过	应急预案	备注
--	-----	-----------	-----	------	------	--------	-----	--------	------	----

## 二、测试内容和结论概述

测试节点硬件配置与软件版本

环境信息	稳定性测试环境 9个机器
------	--------------

CPU	Intel(R) Xeon(R) CPU E5-2680 v4 @ 2.40GHz
内存	256G
网卡	1. Intel Corporation I350 Gigabit Network Connection 2. Intel Corporation 82599EB 10-Gigabit SFI/SFP+
操作系统	发行版: Debian GNU/Linux 9
内核	4.19.87-netease6-1 #2 SMP Mon Sep 7 07:50:31
用途	计算节点
curvefs版本	2.1.0

部署方式

s3	nos
镜像	harbor.cloud.netease.com/curve/curvefs:citest
disk cache	INTEL SSDSC2BB80 800G
metaserver数据	ssd 混合部署
mds	ssd 混合部署
etcd	ssd 混合部署
curveadm版本	0.1.0

三、测试要点

- 1、warmup相关功能、异常、性能测试
- 2、cto相关问题修复

- 
- 3、copysets数据均衡性
  - 4、新版本sdk稳定性和性能

## 四、测试结论

## 五、详细测试数据及监控数据

### 5.1 常规测试

#### 5.1.1 文件系统POSIX 接口

##### 5.1.1.1 pjdtest

已在ci中测试。

##### 5.1.1.2 ltp-fsstress

测试程序: [ltp-full-20220930.tar.bz2](#)

clinet.yaml

```
kind: curvefs

s3.ak: minioadmin

s3.sk: minioadmin

s3.endpoint: 10.182.30.15:9002

s3.bucket_name: chengyi01

container_image: harbor.cloud.netease.com/curve/curvefs:release-2.4-3e38c86

mdsOpt.rpcRetryOpt.addrs: 10.182.2.46:7811,10.182.2.47:7811,10.182.2.48:7811

log_dir: /home/nbs/chengyi01/fs/logs/client

data_dir: /data/chengyi-cache

client.loglevel: 0
```

测试步骤:

```
set -ex

mkdir -p fsstress
pushd fsstress
wget -q -O ltp-full.tgz http://59.111.93.102:8080/qa/ltp-full.tgz //
tar xzf ltp-full.tgz
pushd ltp-full-20091231/testcases/kernel/fs/fsstress
make
BIN=$(readlink -f fsstress)
popd
popd
T=$(mktemp -d -p .)
"$BIN" -d "$T" -l 1 -n 1000 -p 10 -v
echo $?
rm -rf -- "$T"
```

测试结果:

success

## 5.1.2 元数据项& 数据属性

### 5.1.2.1 dbench

dbench

执行命令:

```
sudo dbench -t 600 -D ltp-full-20220930 -c /usr/share/dbench/client.txt 10
```

结果:

Operation	Count	AvgLat	MaxLat
NTCreateX	105740	11.108	1149.330
Close	77763	11.617	1326.275
Rename	4473	500.666	4445.447
Unlink	21257	10.667	1143.985
Qpathinfo	95954	5.922	1405.727
Qfileinfo	16758	0.019	0.138
Qfsinfo	17514	1.216	740.514
Sfileinfo	8567	5.179	684.553
Find	37026	13.418	1599.009
WriteX	52194	0.079	4.078
ReadX	166120	0.196	628.751
LockX	348	0.005	0.015
UnlockX	348	0.002	0.017
Flush	7349	39.296	1008.781

Throughput 5.50159 MB/sec 10 clients 10 procs max\_latency=4445.454 ms

### 5.1.2.2 iozone

测试步骤:

```
iozone -a -n 1g -g 4g -i 0 -i 1 -i 2 -i 3 -i 4 -i 5 -i 8 -f testdir -Rb log.xls
iozone -c -e -s 1024M -r 16K -t 1 -F testfile -i 0 -i 1
iozone -c -e -s 1024M -r 1M -t 1 -F testfile -i 0 -i 1
iozone -c -e -s 10240M -r 1M -t 1 -F testfile -i 0 -i 1
```

测试结果:

```
iozone -a -n 1g -g 4g -i 0 -i 1 -i 2 -i 3 -i 4 -i 5 -i 8 -f testdir -Rb log.xls
```

log.xls

---

```
iozone -c -e -s 1024M -r 16K -t 1 -F testfile -i 0 -i 1
```

Include close in write timing  
Include fsync in write timing  
File size set to 1048576 kB  
Record Size 16 kB  
Command line used: iotest -c -e -s 1024M -r 16K -t 1 -F testfile -i 0 -i 1  
Output is in kBytes/sec  
Time Resolution = 0.000001 seconds.  
Processor cache size set to 1024 kBytes.  
Processor cache line size set to 32 bytes.  
File stride size set to 17 \* record size.  
Throughput test with 1 process  
Each process writes a 1048576 kByte file in 16 kByte records

Children see throughput for 1 initial writers	=	87558.38 kB/sec
Parent sees throughput for 1 initial writers	=	87554.11 kB/sec
Min throughput per process	=	87558.38 kB/sec
Max throughput per process	=	87558.38 kB/sec
Avg throughput per process	=	87558.38 kB/sec
Min xfer	=	1048576.00 kB

Children see throughput for 1 rewriters	=	85795.09 kB/sec
Parent sees throughput for 1 rewriters	=	85790.68 kB/sec
Min throughput per process	=	85795.09 kB/sec
Max throughput per process	=	85795.09 kB/sec
Avg throughput per process	=	85795.09 kB/sec
Min xfer	=	1048576.00 kB

Children see throughput for 1 readers	=	320914.25 kB/sec
Parent sees throughput for 1 readers	=	320818.59 kB/sec
Min throughput per process	=	320914.25 kB/sec
Max throughput per process	=	320914.25 kB/sec
Avg throughput per process	=	320914.25 kB/sec
Min xfer	=	1048576.00 kB

Children see throughput for 1 re-readers	=	983651.06 kB/sec
Parent sees throughput for 1 re-readers	=	983069.33 kB/sec
Min throughput per process	=	983651.06 kB/sec



```
Max throughput per process      = 983651.06 kB/sec
Avg throughput per process      = 983651.06 kB/sec
Min xfer                        = 1048576.00 kB
```

```
iozone -c -e -s 1024M -r 1M -t 1 -F testfile -i 0 -i 1
```

```
Include close in write timing
Include fsync in write timing
File size set to 1048576 kB
Record Size 1024 kB
Command line used: iozone -c -e -s 1024M -r 1M -t 1 -F testfile -i 0 -i 1
Output is in kBytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 kBytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.
Throughput test with 1 process
Each process writes a 1048576 kByte file in 1024 kByte records

Children see throughput for 1 initial writers = 96138.94 kB/sec
Parent sees throughput for 1 initial writers = 96132.17 kB/sec
Min throughput per process = 96138.94 kB/sec
Max throughput per process = 96138.94 kB/sec
Avg throughput per process = 96138.94 kB/sec
Min xfer = 1048576.00 kB

Children see throughput for 1 rewriters = 99591.31 kB/sec
Parent sees throughput for 1 rewriters = 99584.72 kB/sec
Min throughput per process = 99591.31 kB/sec
Max throughput per process = 99591.31 kB/sec
Avg throughput per process = 99591.31 kB/sec
Min xfer = 1048576.00 kB

Children see throughput for 1 readers = 326366.78 kB/sec
Parent sees throughput for 1 readers = 326293.75 kB/sec
Min throughput per process = 326366.78 kB/sec
```

---

Max throughput per process	= 326366.78 kB/sec
Avg throughput per process	= 326366.78 kB/sec
Min xfer	= 1048576.00 kB
Children see throughput for 1 re-readers	= 1567921.62 kB/sec
Parent sees throughput for 1 re-readers	= 1566385.40 kB/sec
Min throughput per process	= 1567921.62 kB/sec
Max throughput per process	= 1567921.62 kB/sec
Avg throughput per process	= 1567921.62 kB/sec

Min xfer = 1048576.00 kB

iozone -c -e -s 10240M -r 1M -t 1 -F testfile -i 0 -i 1

Include close in write timing  
Include fsync in write timing  
File size set to 10485760 kB  
Record Size 1024 kB  
Command line used: iozone -c -e -s 10240M -r 1M -t 1 -F testfile -i 0 -i 1  
Output is in kBytes/sec  
Time Resolution = 0.000001 seconds.  
Processor cache size set to 1024 kBytes.  
Processor cache line size set to 32 bytes.  
File stride size set to 17 \* record size.  
Throughput test with 1 process  
Each process writes a 10485760 kByte file in 1024 kByte records

Children see throughput for 1 initial writers = 97156.06 kB/sec  
Parent sees throughput for 1 initial writers = 97155.48 kB/sec  
Min throughput per process = 97156.06 kB/sec  
Max throughput per process = 97156.06 kB/sec  
Avg throughput per process = 97156.06 kB/sec  
Min xfer = 10485760.00 kB

Children see throughput for 1 rewriters = 96146.71 kB/sec  
Parent sees throughput for 1 rewriters = 96145.96 kB/sec  
Min throughput per process = 96146.71 kB/sec  
Max throughput per process = 96146.71 kB/sec  
Avg throughput per process = 96146.71 kB/sec  
Min xfer = 10485760.00 kB

Children see throughput for 1 readers = 341906.62 kB/sec  
Parent sees throughput for 1 readers = 341898.45 kB/sec  
Min throughput per process = 341906.62 kB/sec  
Max throughput per process = 341906.62 kB/sec

---

Avg throughput per process	= 341906.62 kB/sec
Min xfer	= 10485760.00 kB
Children see throughput for 1 re-readers	= 1456958.62 kB/sec
Parent sees throughput for 1 re-readers	= 1456786.91 kB/sec
Min throughput per process	= 1456958.62 kB/sec
Max throughput per process	= 1456958.62 kB/sec
Avg throughput per process	= 1456958.62 kB/sec

Min xfer

= 10485760.00 kB

#### 5.1.2.3 mdtest

测试步骤:

```
#
for i in 4 8 16;do mpirun --allow-run-as-root -np $i mdtest -z 2 -b 3 -I 10000 -d
/home/chengyi01/test/mnt/mnt-1/mdtest;done

#
mpirun --allow-run-as-root -np 16 mdtest -C -F -L -z 4 -b 10 -I 10000 -d mdtest -w 1024
```

测试结果:

```
for i in 4 8 16;do mpirun --allow-run-as-root -np $i mdtest -z 2 -b 3 -I 10000 -d /home/chengyi01/test/mnt/mnt-1/mdtest;done
```

```
-- started at 12/16/2022 17:37:53 --
```

```
mdtest-1.9.3 was launched with 4 total task(s) on 1 node(s)
```

```
Command line used: mdtest -z 2 -b 3 -I 10000 -d /home/chengyi01/test/mnt/mnt-1/mdtest
```

```
Path: /home/chengyi01/test/mnt/mnt-1
```

```
FS: 10240.0 TiB   Used FS: 0.0%   Inodes: 1024.0 Mi   Used Inodes: 0.0%
```

```
4 tasks, 520000 files/directories
```

```
SUMMARY: (of 1 iterations)
```

Operation	Max	Min	Mean	Std Dev
-----	---	---	----	-----
Directory creation:	659.362	659.362	659.362	0.000
Directory stat :	640.632	640.632	640.632	0.000
Directory removal :	376.895	376.895	376.895	0.000
File creation :	573.503	573.503	573.503	0.000

```
File stat      :      984.815      984.815      984.815      0.000
File read      :      916.783      916.783      916.783      0.000
File removal    :      439.710      439.710      439.710      0.000
Tree creation   :      313.125      313.125      313.125      0.000
Tree removal    :      12.875       12.875       12.875      0.000
```

```
-- finished at 12/16/2022 19:20:39 --
-- started at 12/16/2022 19:20:39 --
```

mdtest-1.9.3 was launched with 8 total task(s) on 1 node(s)  
Command line used: mdtest -z 2 -b 3 -I 10000 -d /home/chengyi01/test/mnt/mnt-1/mdtest  
Path: /home/chengyi01/test/mnt/mnt-1  
FS: 10240.0 TiB Used FS: 0.0% Inodes: 1024.0 Mi Used Inodes: 0.0%

8 tasks, 1040000 files/directories

SUMMARY: (of 1 iterations)

Operation	Max	Min	Mean	Std Dev
-----	---	---	----	-----
Directory creation:	676.693	676.693	676.693	0.000
Directory stat :	1673.279	1673.279	1673.279	0.000
Directory removal :	436.321	436.321	436.321	0.000
File creation :	597.750	597.750	597.750	0.000
File stat :	1698.615	1698.615	1698.615	0.000
File read :	1622.404	1622.404	1622.404	0.000
File removal :	494.133	494.133	494.133	0.000
Tree creation :	310.338	310.338	310.338	0.000
Tree removal :	6.949	6.949	6.949	0.000

```
-- finished at 12/16/2022 22:01:21 --
-- started at 12/16/2022 22:01:21 --
```

mdtest-1.9.3 was launched with 16 total task(s) on 1 node(s)  
Command line used: mdtest -z 2 -b 3 -I 10000 -d /home/chengyi01/test/mnt/mnt-1/mdtest  
Path: /home/chengyi01/test/mnt/mnt-1  
FS: 10240.0 TiB Used FS: 0.0% Inodes: 1024.0 Mi Used Inodes: 0.0%

16 tasks, 2080000 files/directories

---

SUMMARY: (of 1 iterations)

Operation	Max	Min	Mean	Std Dev
-----	---	---	----	-----
Directory creation:	782.754	782.754	782.754	0.000
Directory stat :	2277.546	2277.546	2277.546	0.000
Directory removal :	481.781	481.781	481.781	0.000
File creation :	662.049	662.049	662.049	0.000
File stat :	2276.689	2276.689	2276.689	0.000
File read :	2164.328	2164.328	2164.328	0.000
File removal :	544.474	544.474	544.474	0.000
Tree creation :	315.641	315.641	315.641	0.000
Tree removal :	3.526	3.526	3.526	0.000

---

```
-- finished at 12/17/2022 02:40:09 --
```

#### 5.1.2.4 rename 测试用例集

暂无

#### 5.1.2.5 xfstest

测试步骤:

```
#!/bin/sh -x

set -e

wget http://59.111.93.102:8080/qa/fsync-tester.c
gcc -D_GNU_SOURCE fsync-tester.c -o fsync-tester

./fsync-tester

echo $PATH
whereis lsof
lsof
```

测试结果:

success

#### 5.1.3 数据一致性测试

##### 5.1.3.1 编译项目或者内核

测试步骤:



```
# linux
#!/usr/bin/env bash

set -e

wget -O linux.tar.gz https://mirrors.edge.kernel.org/pub/linux/kernel/v5.x/linux-5.4.tar.xz
sudo apt-get install libelf-dev bc -y
mkdir t
cd t
tar xzf ../linux.tar.gz
cd linux*
make defconfig
make -j`grep -c processor /proc/cpuinfo`
cd ..
if ! rm -rv linux* ; then
    echo "uh oh rm -r failed, it left behind:"
    find .
    exit 1
fi
cd ..
rm -rv t linux*
```

测试结果:

success


#### 5.1.3.2 vdbench读写一致性测试

测试步骤:

```
fsd=fsdl,anchor=/home/nbs/failover/test1,depth=1,width=10,files=10,sizes=(100m,0),shared=yes,openflags=o_direct
fwd=fwdl,fsd=fsdl,threads=10,xfersize=(512,20,4k,20,64k,20,512k,20,1024k,20),fileio=random,fileselect=random,rdp
ct=50
rd=rdl,fwd=fwd*,fwdrate=max,format=restart,elapsed=10000,interval=1

exec : ./vdbench -f profile -jn
```

测试结果:



5.2 异常测试

操作	影响
kill mds 后重启	无影响
kill etcd 后重启	无影响
kill metaserver 后重启	无影响
主etcd掉电	etcd数据如果在系统盘，数据会丢失

主mds掉电	无影响
一个metasever掉电	概率性 io卡顿 0-10s，需要配置数据盘自动挂载
1个etcd\mds\metaserver 网络拔出	无影响
client 网络拔出	网络恢复后恢复
网络延时300ms	性能下降
丢包30%	性能卡顿
丢包10%	性能卡顿
client节点丢包	性能卡顿
metaserver 数据迁出	正常
增加metaserver数据迁入	新增metaserver上无copysets。迁入需要等原故障的metaserver迁出

## 5.3 新增功能测试

### 5.3.1 warmup测试

缓存盘 1.5TiB,

#### 5.3.1.1 cto open

##### 5.3.1.1.1 静态warmup

参考 <http://eq.hz.netease.com/#/useCaseManag/list?projectId=1155&moduleid=9870838> 中的 fs文件系统/2.4.0版本自测用例/预热数据

##### 5.3.1.1.2 同时有读写时warmup

###### 5.3.1.1.2.1 缓存盘容量不足时

调整缓存盘的容量（fdisk），制造缓存盘容量不足。

缓存盘大小保持一致。

###### 5.3.1.1.2.1.1 大文件（根据缓存盘容量）并发操作

操作	结论
挂卸载 fuse	md5一致
其他文件并发读写	md5一致
单metaserver异常（kill）	md5一致
多挂载，不共用缓存盘，并发warmup同一文件	md5一致
多挂载，共用缓存盘，并发warmup同一文件	md5一致
多挂载，共用缓存盘，并发warmup不同文件	md5一致

5.3.1.1.2.1.2 大规模目录（1000w+）

操作	结论
挂卸载 fuse	缓存盘内s3对象文件数量不一致
其他目录并发读写	缓存盘内s3对象文件数量不一致
单metaserver异常（kill）	缓存盘内s3对象文件数量不一致
多挂载，不共用缓存盘，并发warmup同一目录	缓存盘内s3对象文件数量不一致
多挂载，共用缓存盘，并发warmup同一目录	缓存盘内s3对象文件数量不一致
多挂载，共用缓存盘，并发warmup不同目录	缓存盘内s3对象文件数量不一致

5.3.1.1.2.2 缓存盘容量足够时

5.3.1.1.2.2.1 大文件（根据缓存盘容量）并发操作

操作	结论
挂卸载 fuse	md5一致
其他文件并发读写	md5一致
单metaserver异常（kill）	md5一致
多挂载，不共用缓存盘，并发warmup同一文件	md5一致
多挂载，共用缓存盘，并发warmup同一文件	md5一致

---

多挂载，共用缓存盘，并发warmup不同文件	md5一致
------------------------	-------

5.3.1.1.2.2.2 大规模目录（1000w+）

操作	结论
挂卸载 fuse	缓存盘内s3对象文件数量一致
其他目录并发读写	缓存盘内s3对象文件数量一致
单metaserver异常（kill）	缓存盘内s3对象文件数量一致
多挂载，不共用缓存盘，并发warmup同一目录	缓存盘内s3对象文件数量一致
多挂载，共用缓存盘，并发warmup同一目录	缓存盘内s3对象文件数量一致
多挂载，共用缓存盘，并发warmup不同目录	缓存盘内s3对象文件数量一致

5.4 回归测试

- <https://github.com/opencurve/curve/issues/1833>
- <https://github.com/opencurve/curve/issues/1841>
- <https://github.com/opencurve/curve/issues/1842>
- <https://github.com/opencurve/curve/issues/1881>
- <https://github.com/opencurve/curve/issues/1851>
- <https://github.com/opencurve/curve/issues/1854>

5.5 2.3升级2.4 测试

1. 集群先升级2.4
2. client再测试升级2.4